



Parallel Computing Aspects

Relevant in the Context of Chrono/Computational Dynamics



Discussion Points

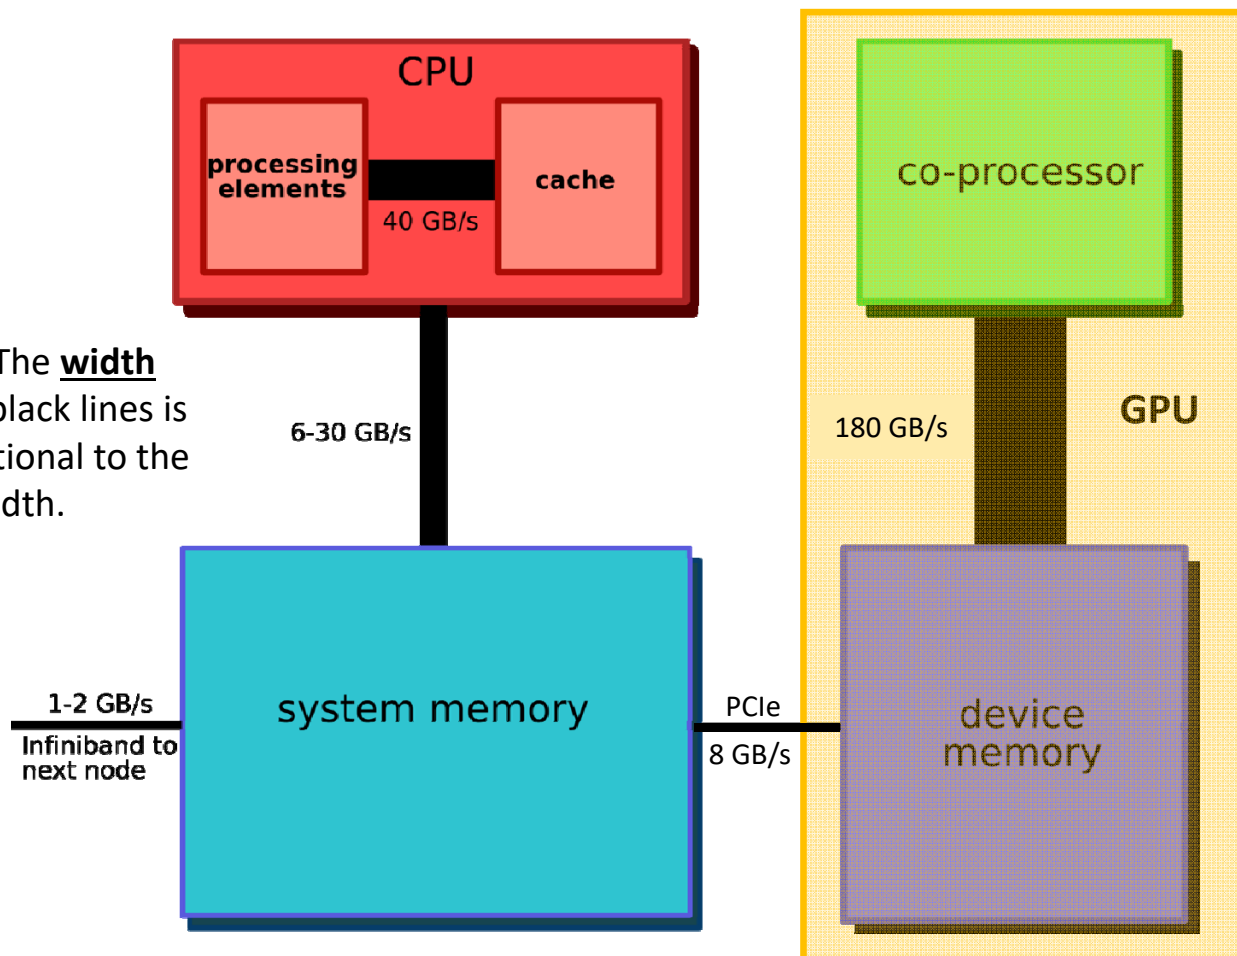
- Computing speed issues, general remarks
- GPU computing and relevance to Chrono
- Multi-core CPU computing and relevance to Chrono
- Distributed-memory computing and relevance to Chrono

The Price of 1 Gflop/second

- How much would you have to pay to perform 1,000,000,000 operations in one second?
 - 1961:
 - Combine 17 million IBM-1620 computers
 - At \$64K apiece, when adjusted for inflation, this would cost \$8.3 trillion
 - 2000:
 - About \$1,000
 - 2015 (January):
 - 8 cents

Commodity CPU-GPU System

NOTE: The **width** of the black lines is proportional to the bandwidth.



Beyond “commodity”: What’s a Powerful Workstation Today?



- Price tag: \$50K

- IBM Power S822LC

- Two Power8 CPUs
- Four to eight Pascal P100 GPUs
- 0.5 TB of system memory
- NVLink (instead of PCIe bus)

SPECIFICATIONS, Tesla P100

GPU Architecture	NVIDIA Pascal
NVIDIA CUDA® Cores	3584
Double-Precision Performance	5.3 TeraFLOPS
Single-Precision Performance	10.6 TeraFLOPS
Half-Precision Performance	21.2 TeraFLOPS
GPU Memory	16 GB CoWoS HBM2
Memory Bandwidth	732 GB/s
Interconnect	NVIDIA NVLink
Max Power Consumption	300 W
ECC	Native support with no capacity or performance overhead
Thermal Solution	Passive
Form Factor	SXM2
Compute APIs	NVIDIA CUDA, DirectCompute, OpenCL™, OpenACC



Today's Workstation Is Yesterday's Supercomputer

- Today one can get about 25-45 TFlops out of a powerful workstation
 - 25 thousands billion operations per second
- 2006: Sandia National Lab Supercomputer
 - 43.5 TFlops
 - Hardware provider: Cray
 - 9th fastest supercomputer in the world
 - Required of the order of tens of KW of power

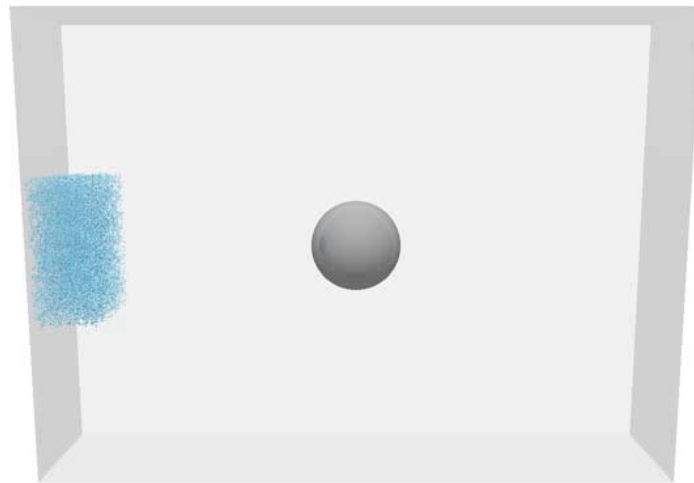
Computational Dynamics: Putting Things in Perspective



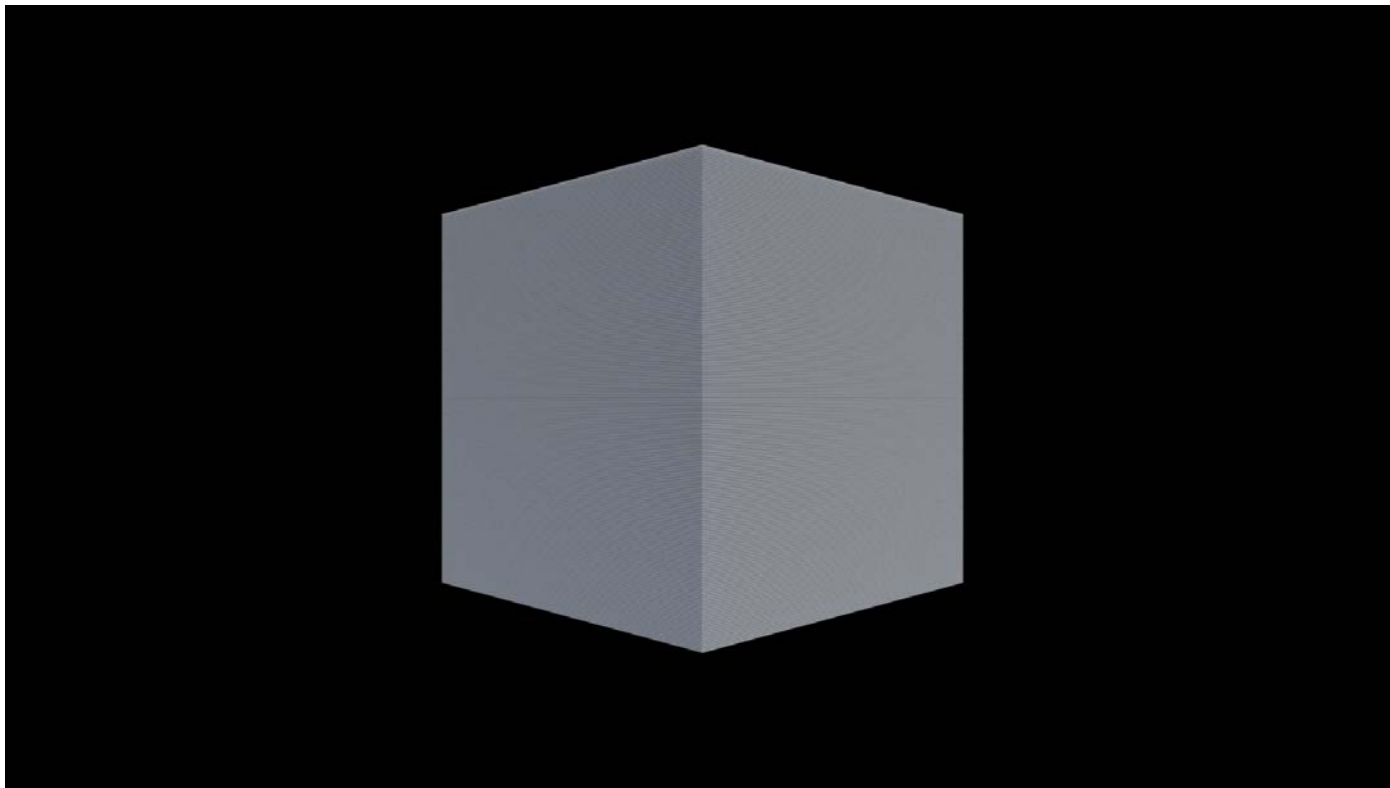
- We can count on lots of compute power
 - Bottom line
 - More accurate models
 - Bridging of space and/or time multiple scales
 - Emerging hardware takes us half way there. Also need
 - Modeling methods + suitable solution algorithms
 - Targeted software, platform specific

Progress over Time

Year: 2009 – 1 million bodies [GPU Granular Dynamics]



Year: 2015 – 10 million bodies



GPU Tank Wave Simulation

- 2009 Implementation

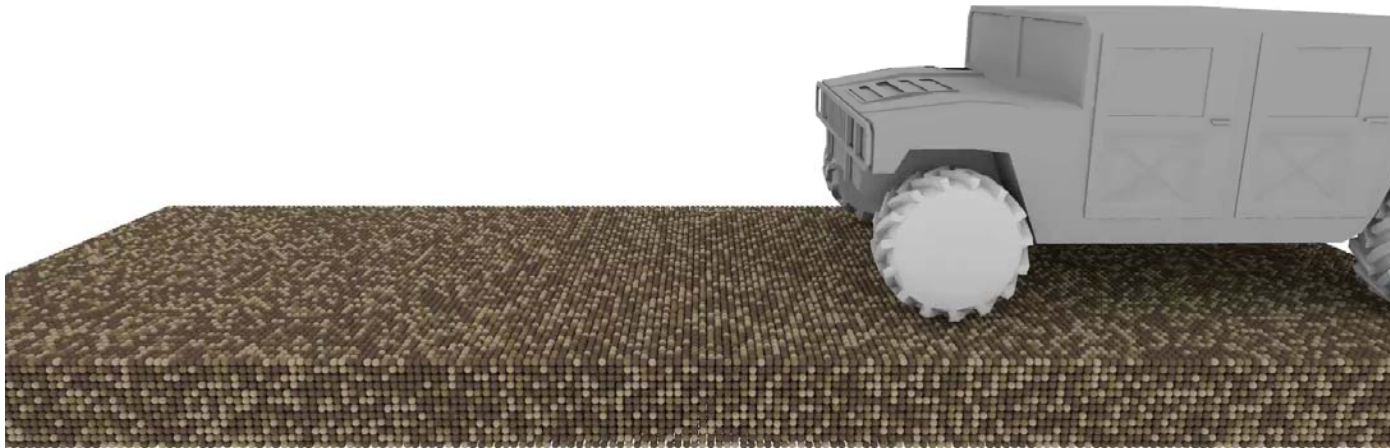
- 1 million rigid frictionless spheres
- Integration time step: 0.01s
- 24 sec per step
- Hardware used: GPU
 - NVIDIA Tesla C1060
- Velocity based complementarity
- 20 seconds long simulation
- Simulation Time: ~2 days

- 2015 Implementation

- 10,648,000 rigid spheres
- Integration time step: 0.00025s
- 1 second per step
- Hardware Use: GPU
 - NVIDIA Tesla K40x
- Position based complementarity
- 10 seconds long simulation
- Simulation Time: 11 hours

2015 Simulation: given the number of bodies, about 20 times faster

Vehicle on Deformable Terrain. Year: 2012



Vehicle on Deformable Terrain. Year: 2015



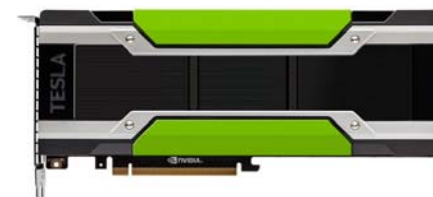
Vehicle on Deformable Terrain

- 2012
 - 300k rigid spheres
 - Length of simulation: 15 seconds
 - Hardware used: CPU (Intel)
 - Multicore, based on OpenMP
 - Integration time step: 0.001s
 - Velocity Based Complementarity
 - 17 seconds per time step
 - Simulation time: ~2.5 days
- 2015
 - ~1.5 million rigid spheres
 - Length of simulation: 15 seconds
 - Hardware: GPU (NVIDIA)
 - Tesla K40X
 - Integration time step: 0.0005s
 - Position Based Dynamics
 - 0.3 seconds per time step
 - Simulation time: ~2.5 hours

2015 Simulation: although 5X more bodies, runs about 25 times faster

Parallel Computing on the GPU

- NVIDIA GPU Computing Architecture
 - Via a separate HW interface
 - In laptops, desktops, workstations, servers
- Multithreaded SIMT model uses application data parallelism and thread parallelism
- Programmable in C with CUDA tools
 - “Extended C”



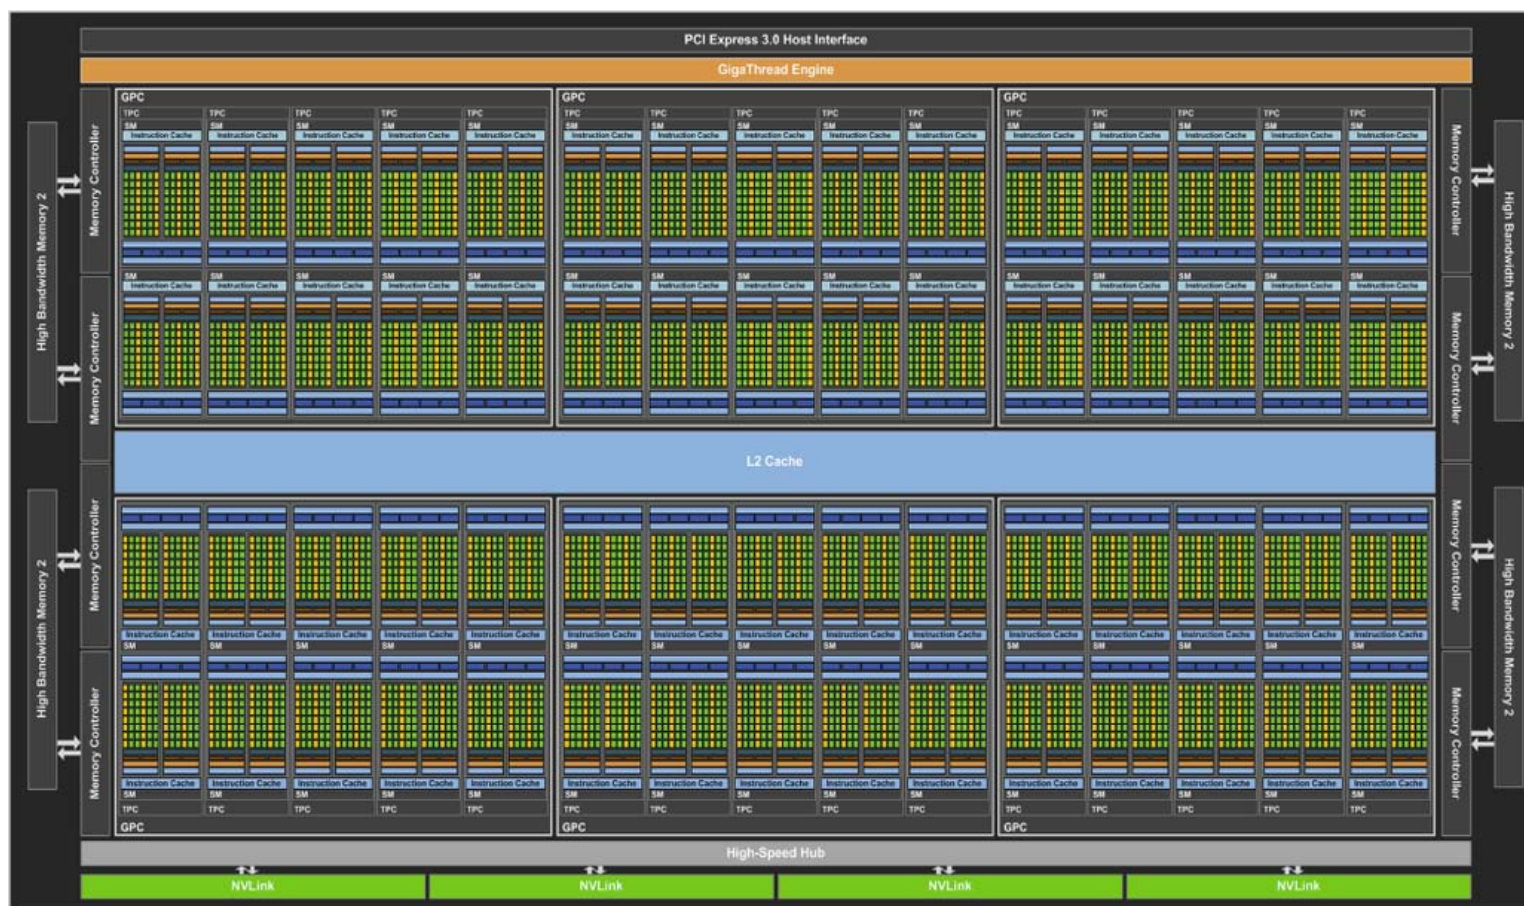
Pascal P100



Kepler K20X

L x W x H: 18 x 8 x 12 inches

Latest NVIDIA GPU Architecture [Tesla P100]



Moore's Law At Work



Tesla Products	Tesla K40	Tesla M40	Tesla P100
GPU	GK110 (Kepler)	GM200 (Maxwell)	GP100 (Pascal)
SMs	15	24	56
TPCs	15	24	28
FP32 CUDA Cores / SM	192	128	64
FP32 CUDA Cores / GPU	2880	3072	3584
FP64 CUDA Cores / SM	64	4	32
FP64 CUDA Cores / GPU	960	96	1792
Base Clock	745 MHz	948 MHz	1328 MHz
GPU Boost Clock	810/875 MHz	1114 MHz	1480 MHz
Peak FP32 GFLOPs ¹	5040	6840	10600
Peak FP64 GFLOPs ¹	1680	210	5300
Texture Units	240	192	224
Memory Interface	384-bit GDDR5	384-bit GDDR5	4096-bit HBM2
Memory Size	Up to 12 GB	Up to 24 GB	16 GB
L2 Cache Size	1536 KB	3072 KB	4096 KB
Register File Size / SM	256 KB	256 KB	256 KB
Register File Size / GPU	3840 KB	6144 KB	14336 KB
TDP	235 Watts	250 Watts	300 Watts
Transistors	7.1 billion	8 billion	15.3 billion
GPU Die Size	551 mm ²	601 mm ²	610 mm ²
Manufacturing Process	28-nm	28-nm	16-nm FinFET

¹ The GFLOPs in this chart are based on GPU Boost Clocks.

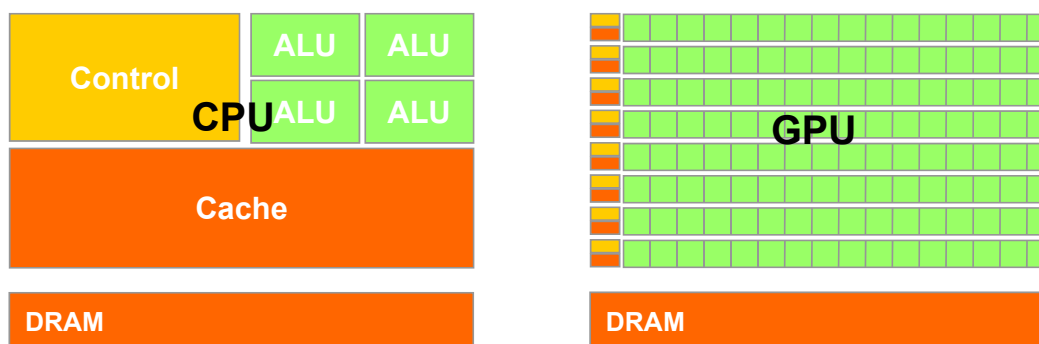


Data vs. Task Parallelism

- Data parallelism
 - You have a large amount of data elements and each data element needs to be processed to produce a result
 - When this processing can be done in parallel, we have data parallelism
 - Example:
 - Filtering a picture; i.e., processing each pixel out of a million in a snapshot
- Task parallelism
 - You have a collection of tasks that need to be completed
 - If these tasks can be performed in parallel you are faced with a task parallel job
 - Examples:
 - Microwave a soup, make a salad, boil pasta, bake a cake
 - All of the above can happen at the same time

Why is the GPU Fast?

- The GPU is **specialized** for compute-intensive, highly data parallel computation (owing to its graphics rendering origin)
 - More transistors devoted to **data processing** rather than data caching and control flow
 - Where are GPUs good: **high arithmetic intensity** (the ratio between arithmetic operations and memory operations)



- Video gaming industry exerts strong **economic pressure** that fuels constant innovation in GPU computing

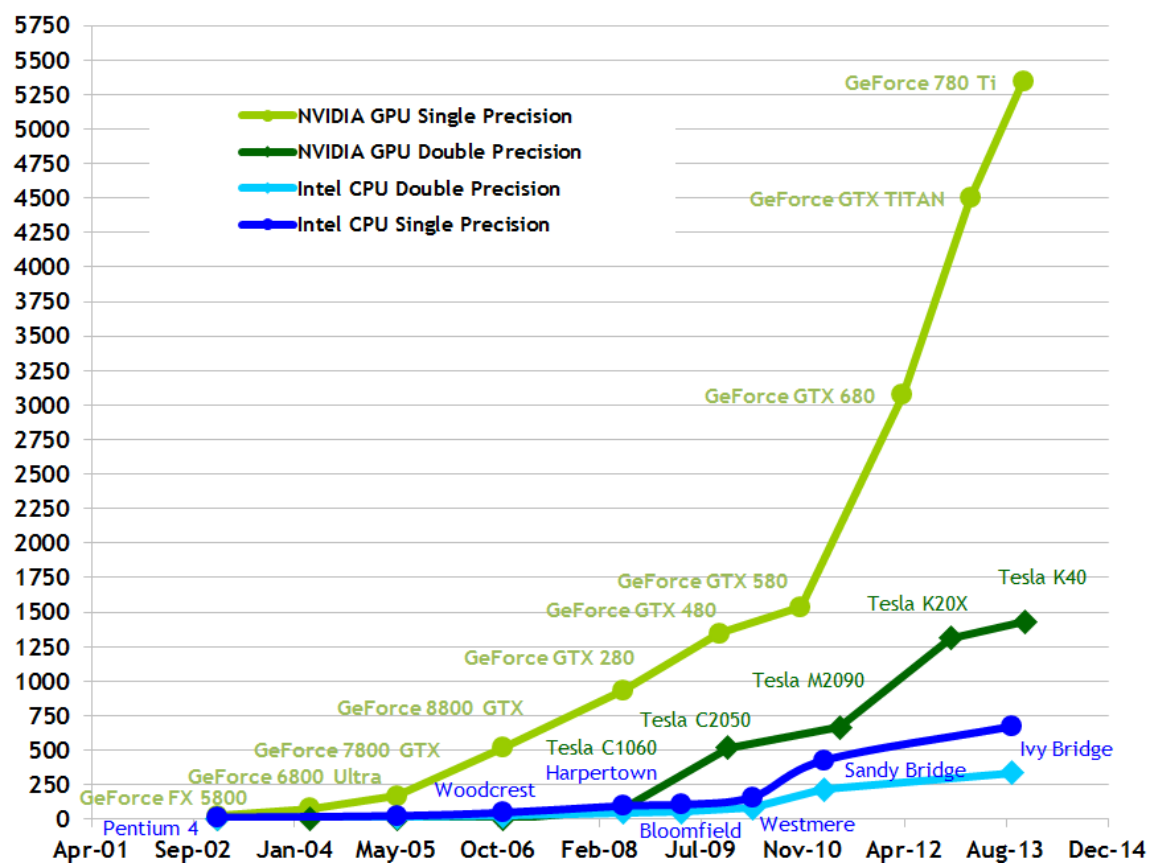
When Are GPUs Good?

- Ideally suited for data-parallel computing (SIMD)
- Moreover, you want to have high arithmetic intensity
 - Arithmetic intensity: ratio of arithmetic operations to memory operations
- You are off to a good start with GPU computing if you can do this...
 - GET THE DATA ON THE GPU AND KEEP IT THERE
 - GIVE THE GPU ENOUGH WORK TO DO
 - FOCUS ON DATA REUSE WITHIN THE GPU TO AVOID MEMORY BANDWIDTH LIMITATIONS

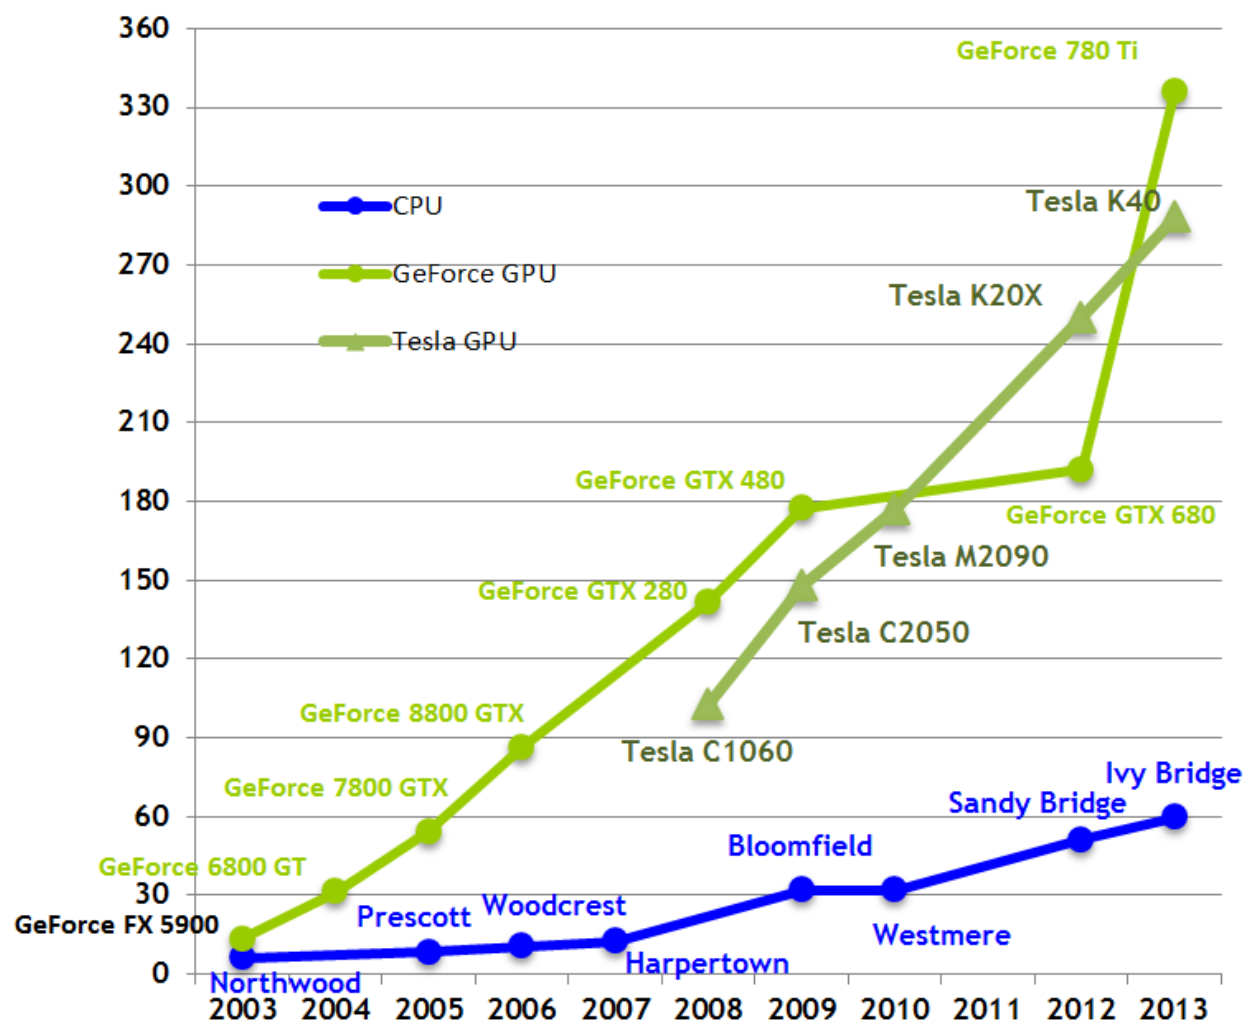
Very Important Point

- Almost 100% of our applications bound by memory speed
 - Most often, the cores (or SPs) idle waiting for data
 - What's important is how fast you can move data
 - You want high memory bandwidth

Theoretical GFLOP/s



Theoretical GB/s



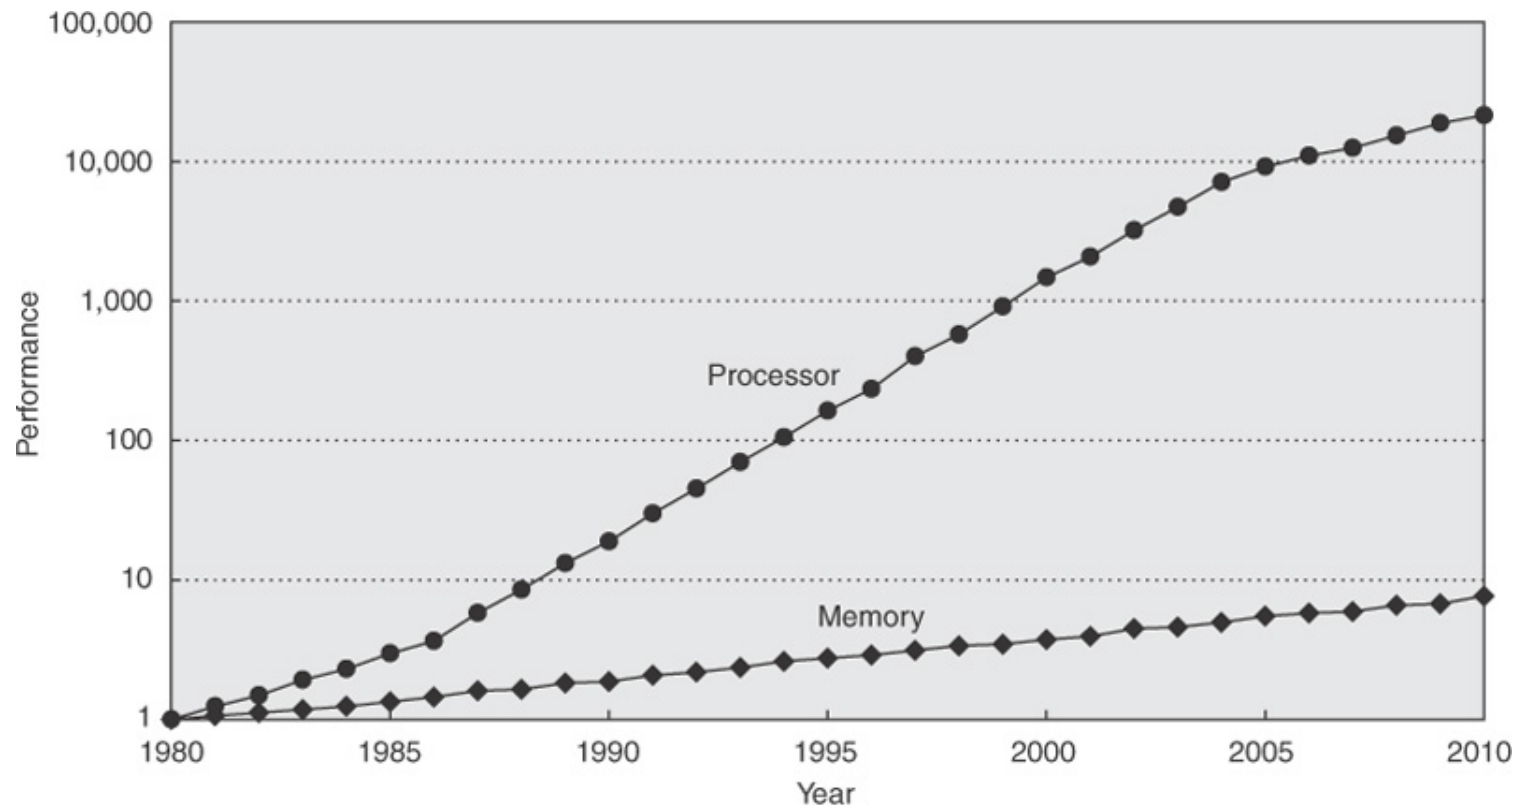
Memory Speed:

Widening of the Processor-DRAM Performance Gap

- The processor: So fast it left the memory behind
 - A system (CPU-Memory duo) can't move as fast as you'd like (based on CPU top speeds) with a sluggish memory
- Plot on next slide shows on a *log* scale the increasing gap between CPU and memory
- The memory baseline: 64 KB DRAM in 1980
- Memory speed increasing at a rate of approx 1.07/year
 - However, processors improved
 - 1.25/year (1980-1986)
 - 1.52/year (1986-2004)
 - 1.20/year (2004-2010)

Memory Speed:

Widening of the Processor-DRAM Performance Gap



© 2007 Elsevier, Inc. All rights reserved.

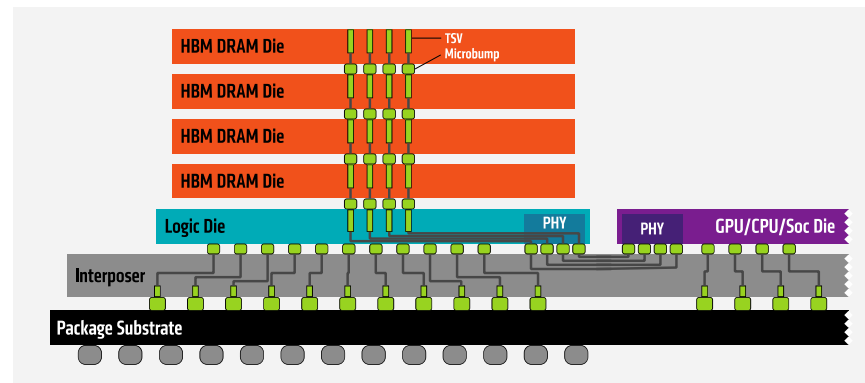
Courtesy of Elsevier, Computer Architecture, Hennessey and Patterson, fourth edition

3D Stacked Memory

[immediate future looks bright]



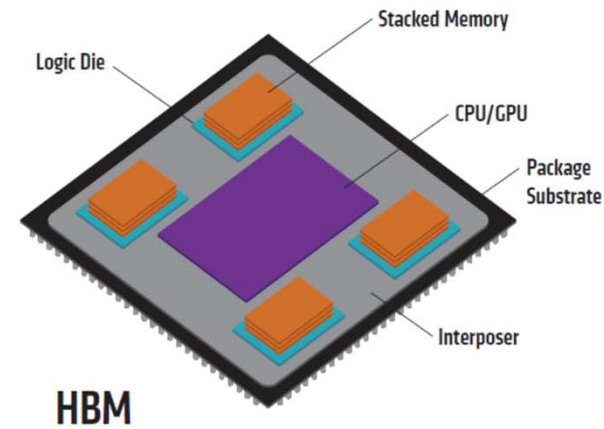
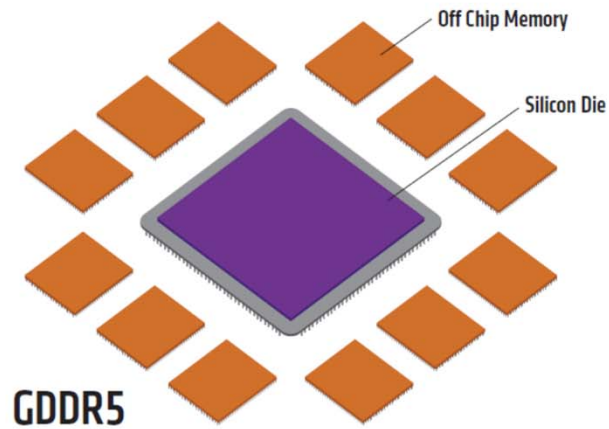
- SK Hynix's High Bandwidth Memory (HBM)
 - Developed by AMD and SK Hynix



- 1st Generation (HBM1) introduced in AMD Fiji GPUs
 - 1GB & 128GB/s per stack
 - AMD Radeon R9 Fury X: had four stacks → 4GB & 512GB/s
- 2nd Generation (HBM2) will be used in NVIDIA Pascal and AMD Arctic Island GPUs
 - 2 GB & 256GB/s bandwidth per stack
 - NVIDIA Pascal: close to 1TB/s memory bandwidth

3D Stacked Memory

[immediate future looks bright]

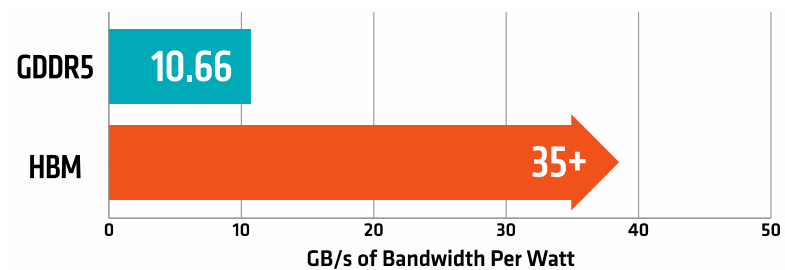
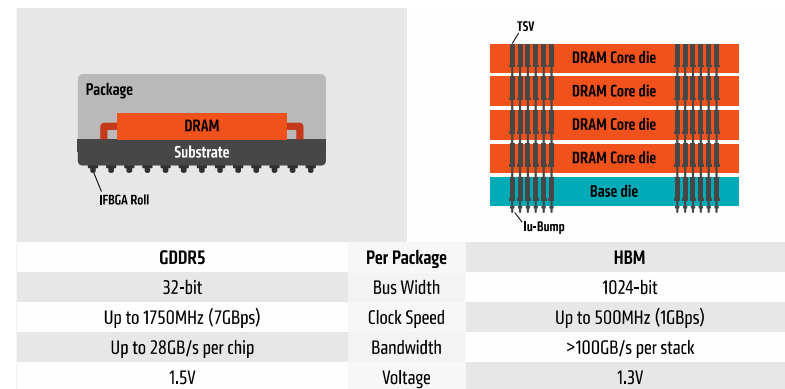
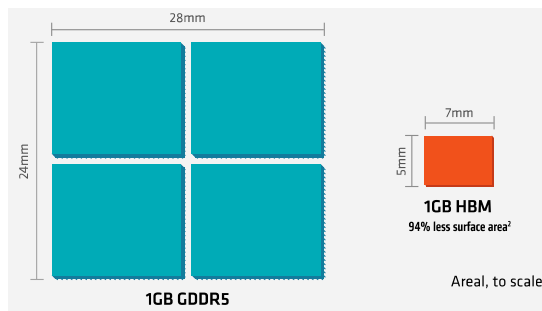


This is a big deal.

3D Stacked Memory

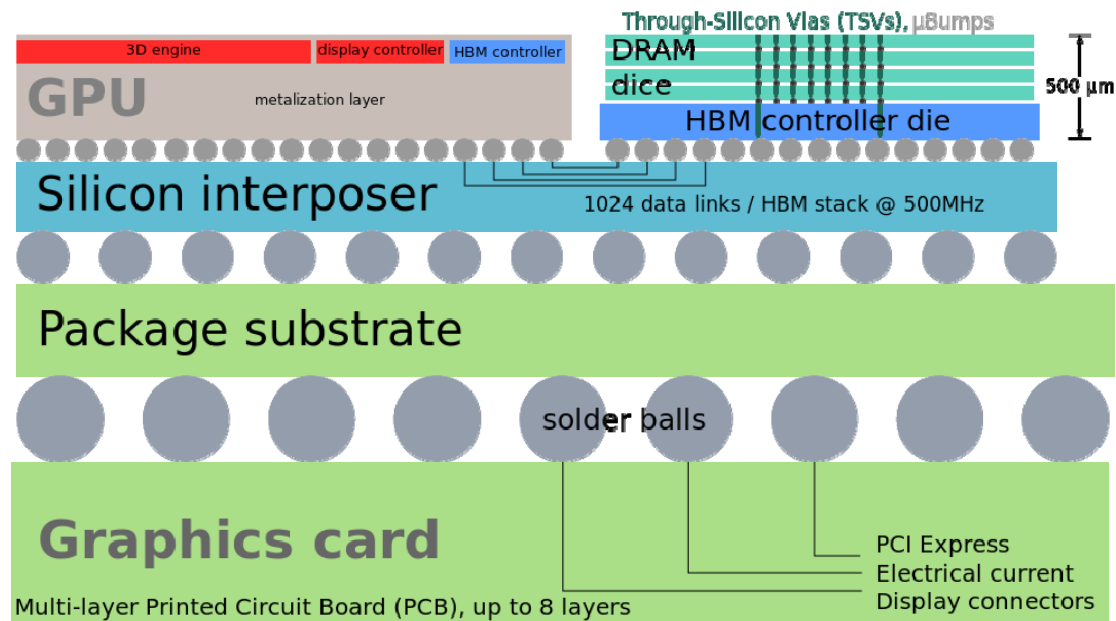


- More power efficient
- Shorter distances
 - Less power wasted moving data
- Smaller memory footprint
 - More memory can be packed into space



GPU w/ HBM


[cut-through]



The Advanced Computing Landscape

Supercomputing (distributed memory)	Group of nodes communicate through fast interconnect	MPI, Charm++, Chapel
Multi-socket Parallel Computing	Group of CPUs operate together on the same node	OpenMP, MPI
Acceleration (GPU/Phi)	Compute devices accelerating parallel computation on one node	CUDA, OpenCL + OpenMP, MPI
Multi-Core Parallel Computing	Communication through shared caches	OpenMP, TBB, pthreads
Vectorization	Higher operation throughput via special/fat registers	AVX, SSE
Pipelining	Sequence of instruction sharing functional units	Assembly
Superscalar	Non-sequence instructions sharing functional units	Assembly

We have full control → 

We have little to no control → 

New Topic:

Looking beyond GPU computing

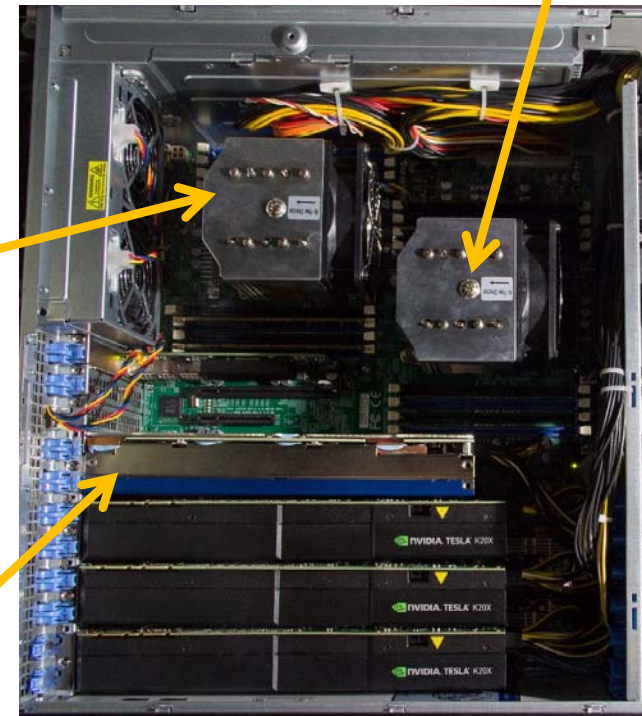


- Much more to parallel computing than GPU acceleration
- Intel: shared memory parallelism
 - “Parallel multicore” computing
- Where are the multiple cores?

10 cores (20 virtual)
E5-2690 v2 Ivy Bridge-EP
3.0GHz 25MB L3 Cache

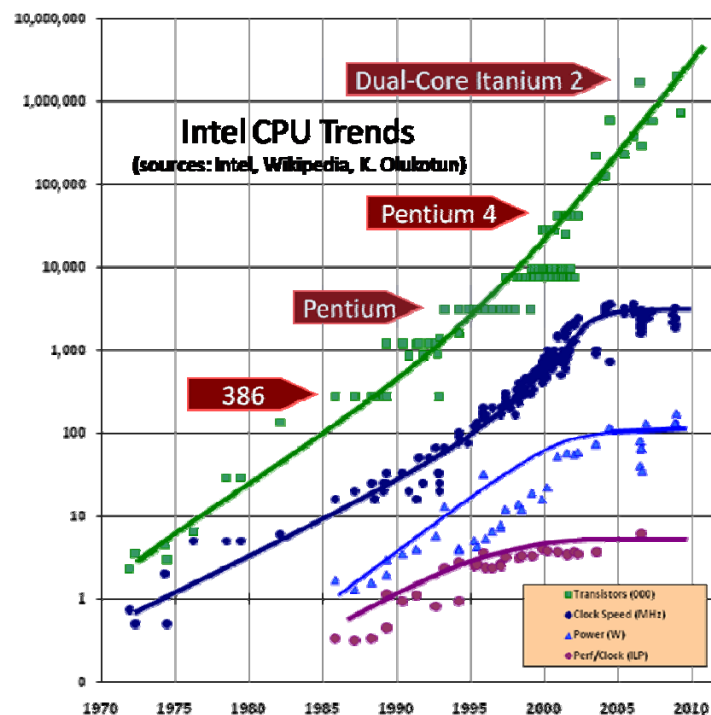
10 cores (20 virtual)
E5-2690 v2 Ivy Bridge-EP
3.0GHz 25MB L3 Cache

60 cores
Intel® Xeon Phi™ 5110P (8GB)



Moore's Law

- Number of transistors per unit area has been steadily going up
- ILP and Clock Speed have stagnated



Transistor Densities Still Going Up For Now

- Consequence: lots of cores in one chip
 - CPU Cores: 18 today, probably **32** by 2018
 - Intel Xeon Phi : 61 today, very likely close to **200** in 2017
- This trend leveraged in Chrono by use of OpenMP

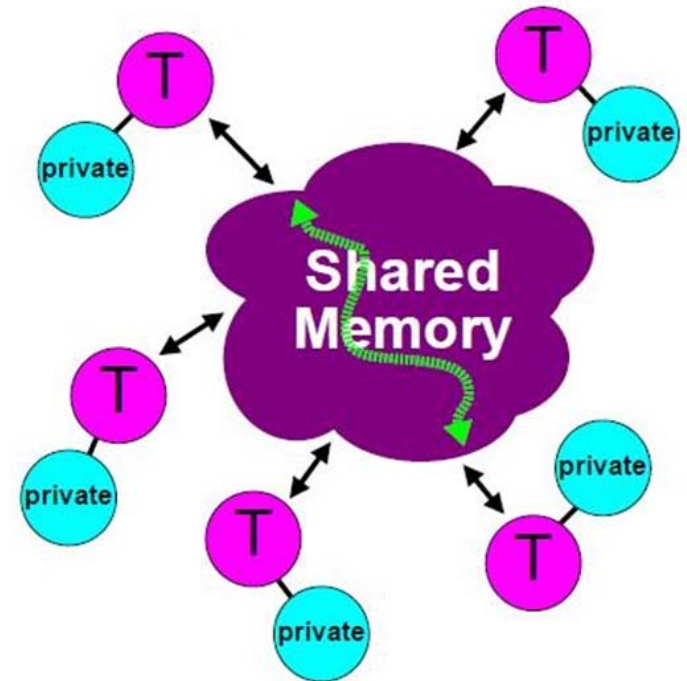
What Is OpenMP?

- Portable, shared-memory threading API
 - Bindings: Fortran, C, and C++
 - Multi-vendor support for both Linux and Windows
- Standardizes task & loop-level parallelism
- Very good at coarse-grained (task) parallelism
- Combines serial and parallel code in single source
- Standardizes ~ 25 years of compiler-directed threading experience
- Current spec is OpenMP 4.0
 - Released in 2013
 - <http://www.openmp.org>
 - More than 300 Pages

Shared Memory Multi-Processing [SMMP]



- Threads have access to large pool of shared memory
- Threads can have private data
 - Not accessible by other threads
- Data transfer/access transparent to programmer
- Synchronization is implicit but can be made explicit as well



OpenMP: Directives-based API

- Most OpenMP constructs are compiler **directives** or **pragmas**

- For C and C++, the pragmas take the form:

`#pragma omp construct [clause [clause]...]`

- For Fortran, the directives take one of the forms:

`C$OMP construct [clause [clause]...]`

`!$OMP construct [clause [clause]...]`

`*$OMP construct [clause [clause]...]`

Work Sharing

- **Work sharing**: term used in OpenMP to describe distribution of work across threads
- Three primary avenues for work sharing in OpenMP:
 - `omp for` construct
 - `omp sections` construct
 - `omp task` construct

Each of them **automatically** divides work among threads

Intel Roadmap: The Full Story

- 2013 – 22 nm Tick: Ivy Bridge – Tock: Haswell
 - 2015 – 14 nm Tick: Broadwell – Tock: Skylake
 - 2016 – 14nm “Refresh” Kabylake
 - 2017 – 10 nm Tick: Cannonlake (delayed to 2nd Half 2017)
 - 2019 – 7 nm
 - 2021 – 5 nm
 - 2023 – ??? (carbon nanotubes?)
-
- Happening now: Moore’s law moving from 18-24 month cycle to 24-30 month cycle for the first time in 50 years
 - In 5 years, end of Moore’s law: can’t increase transistor density on unit area after 2021

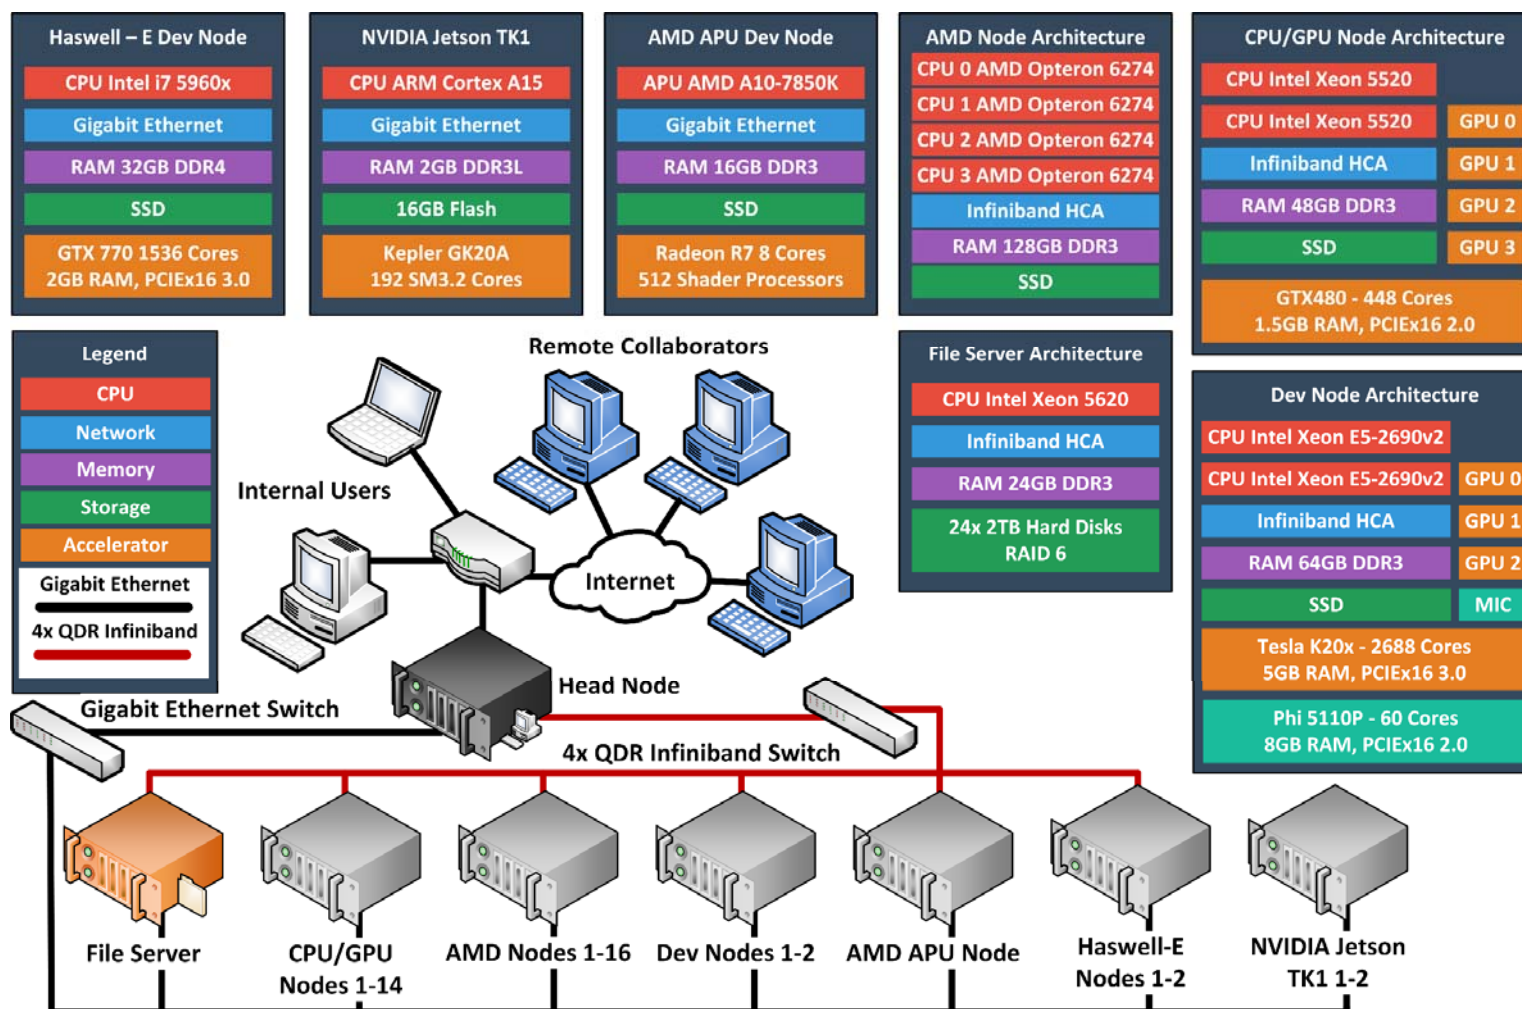
Why's OpenMP Faltering?

- Cache Coherence
 - Many cores see the same memory space and write all over the place
 - Memory manager unit (MMU): the watchdog making sure Thread X running on Core X does not work with data that is stale (since modified by Thread Y running on Core Y)
- NUMA: Non-uniform memory access
 - Some memory banks closer to some cores than to some others
 - Long wait time if a Thread keeps looking for data stored in a far-away memory bank

Scaling, with Lots of Cores: Via Distributed Memory

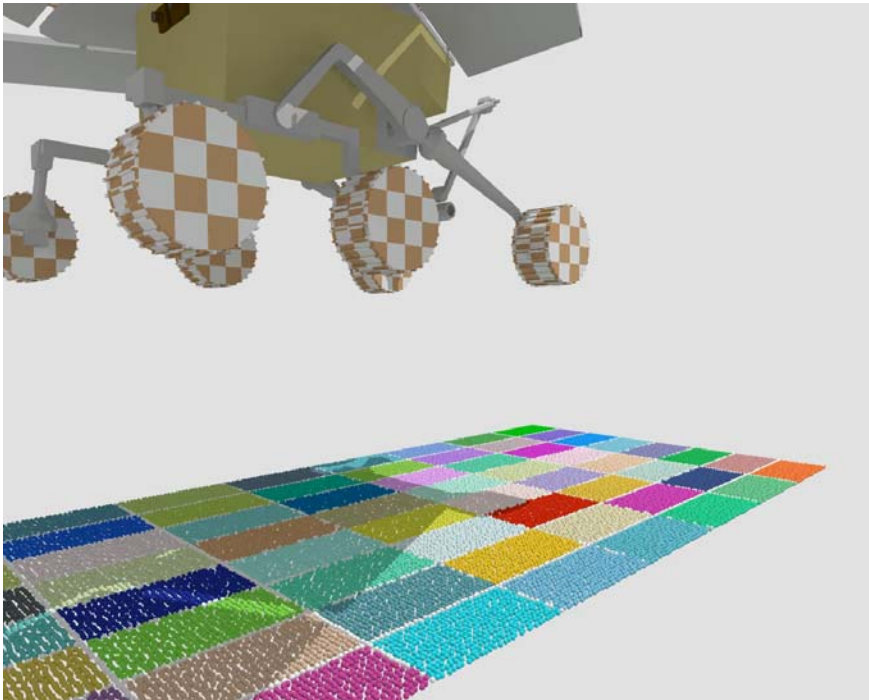
- To scale, a different parallel programming paradigm needed: [distributed memory](#)
- Pros, distributed memory
 - Eliminates cache coherence issues
 - Can attack very large problems – lots of memory available to user
- Cons, distributed memory
 - Calls for major code re-write
 - Very often high data access latencies
 - 1000X higher than accessing memory on a workstation





EULER – Lab's heterogeneous CPU/GPU Cluster.

Distributed Memory Simulation in Chrono



Distributed Memory, Good for the Long Run Though

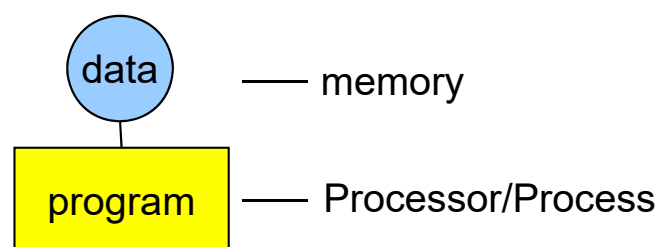
- Five to six years from now, it's not clear what will replace Moore's law
- No technology yet to continue the steady increase in core count we have experienced
 - Can't improve anymore speeds by use of more cores on one workstation
- Distributed memory is the path towards running by drawing on multiple workstations
 - Called "nodes"

Distributed Memory Solution

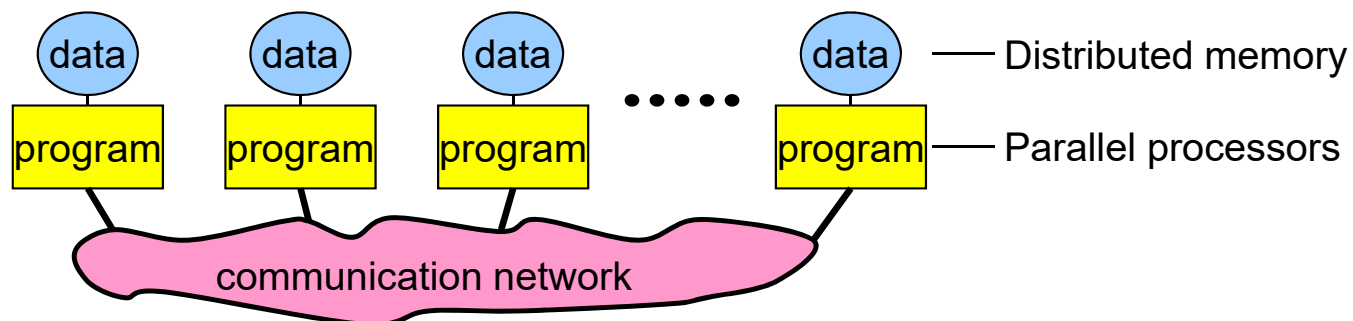
- Distributed Memory Solution = Supercomputer
- Tianhe, IBM BlueGene, Cray, Linux cluster, etc.
- Almost all of these machines rely on the Message Passing Interface (MPI)
- MPI not quite friendly to large scale multibody dynamics simulation

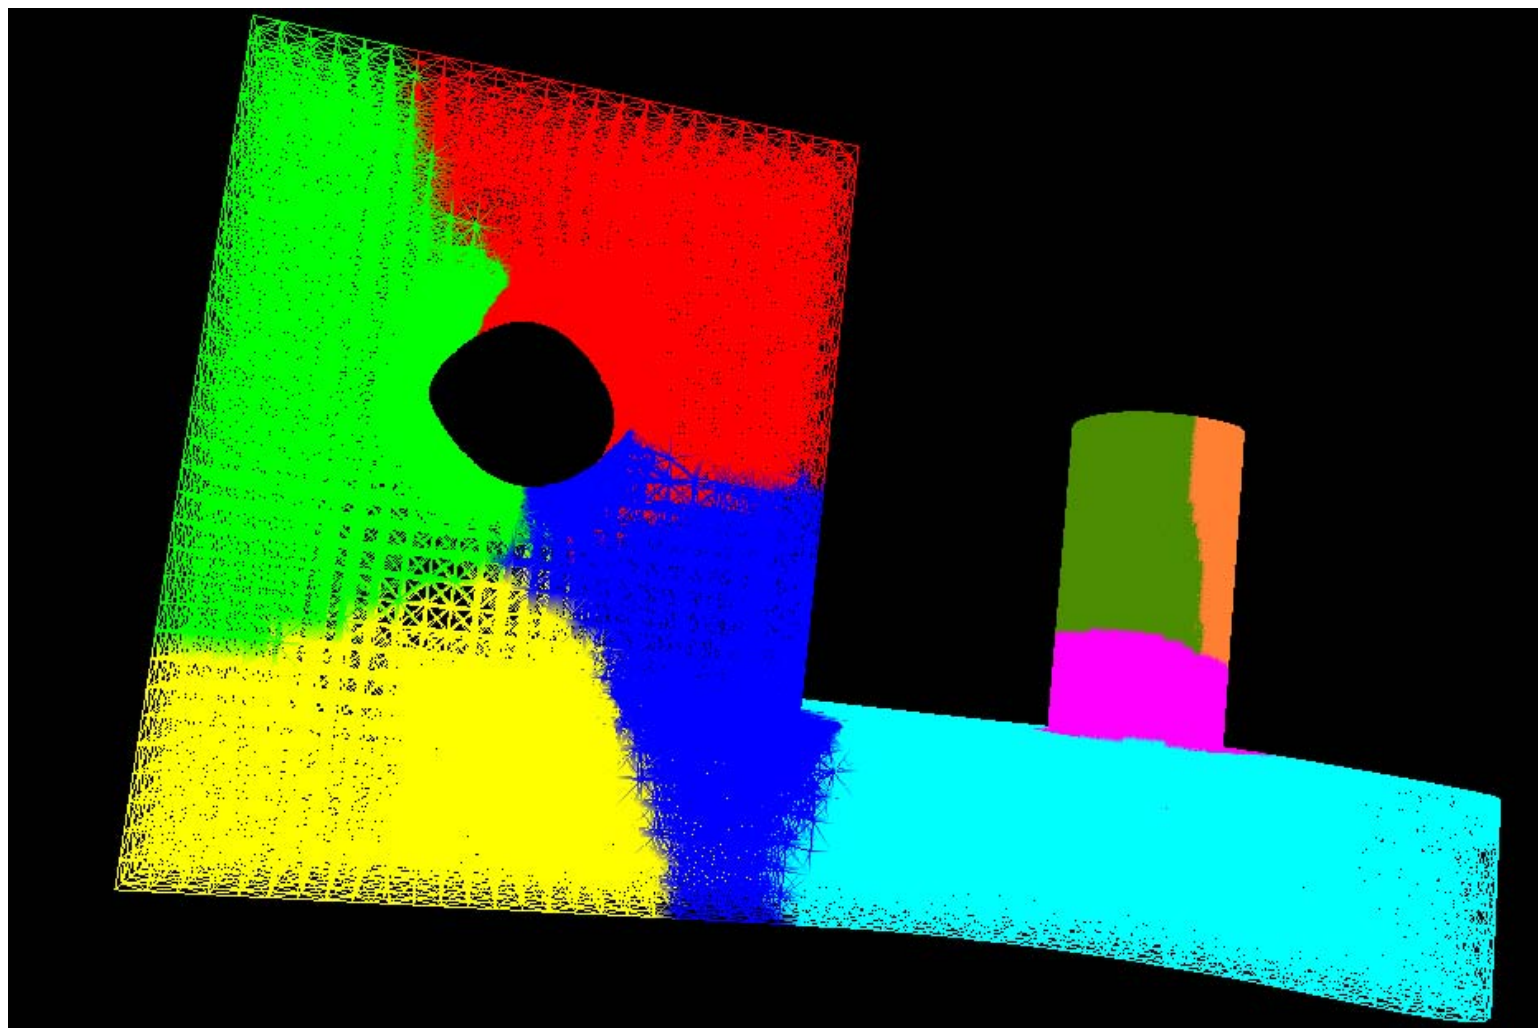
The Message-Passing Programming Paradigm

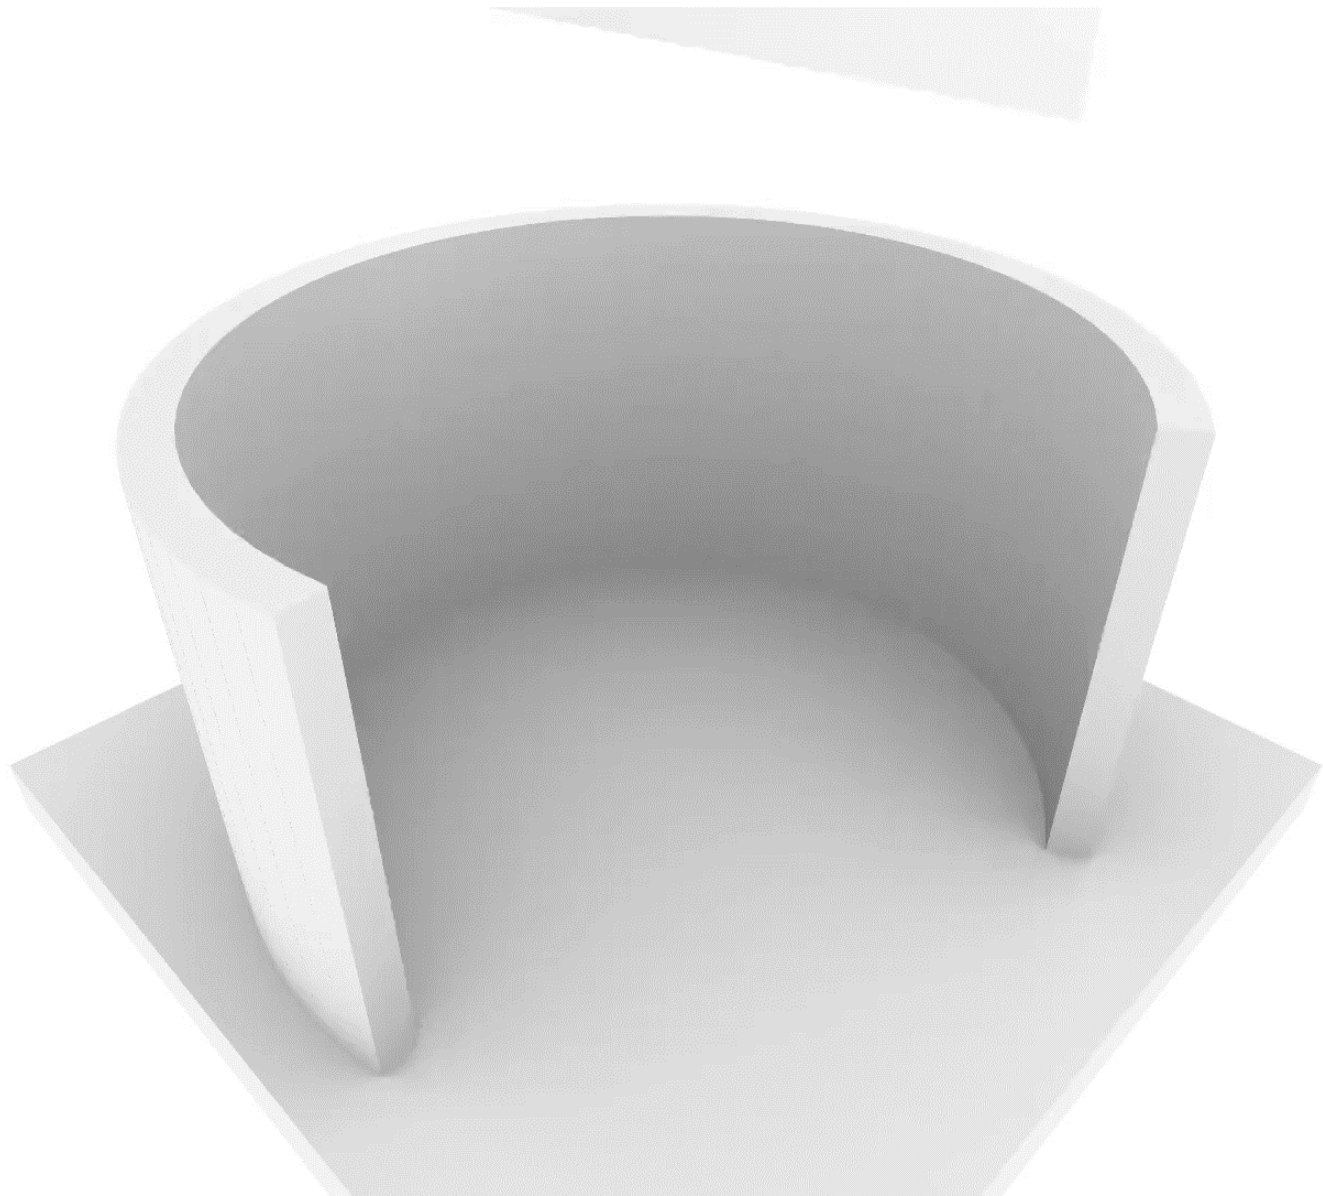
- Sequential Programming Paradigm



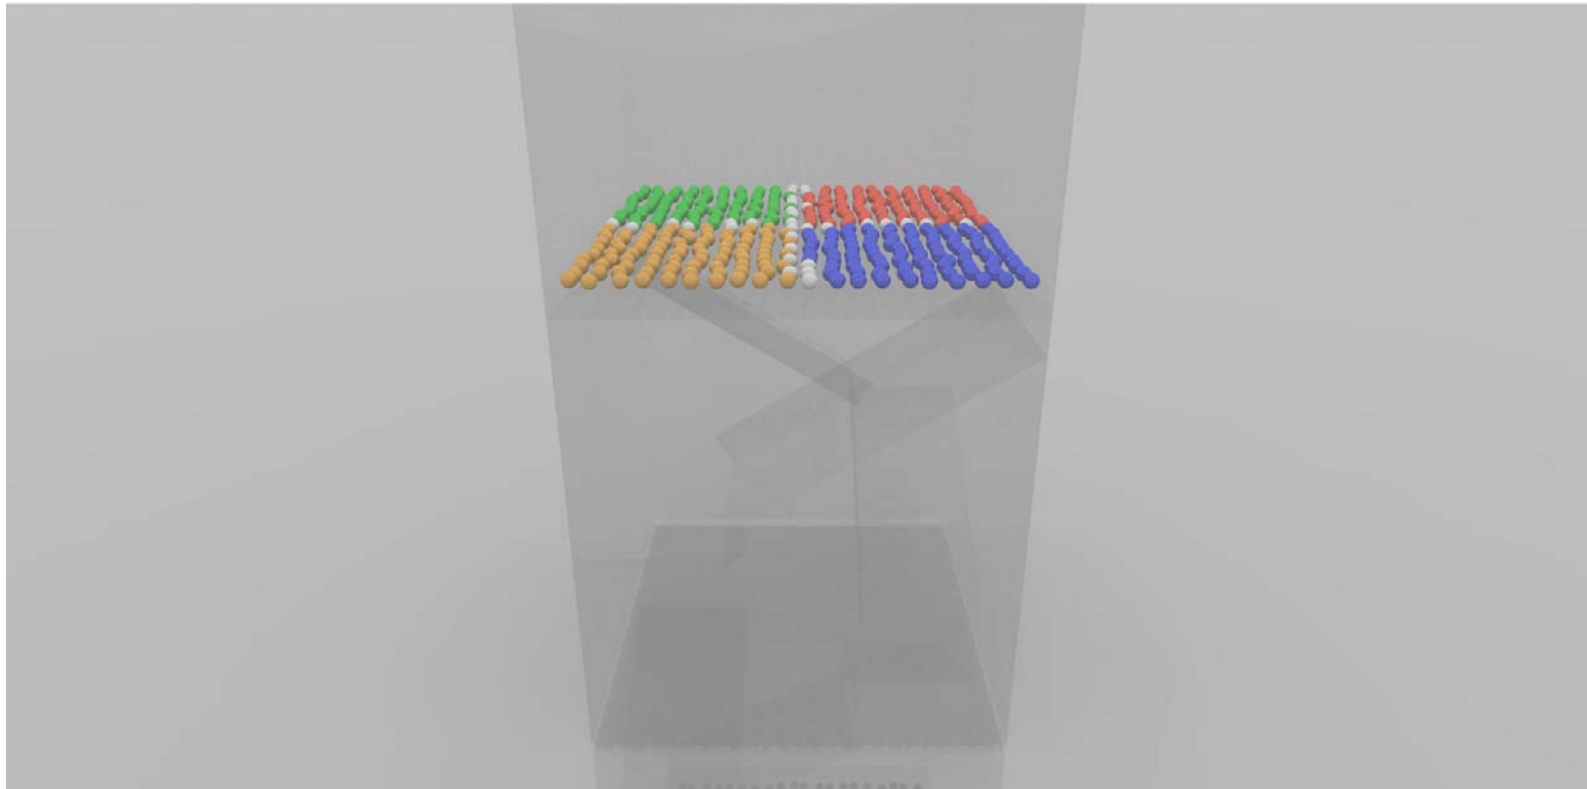
-
- Message-Passing Programming Paradigm





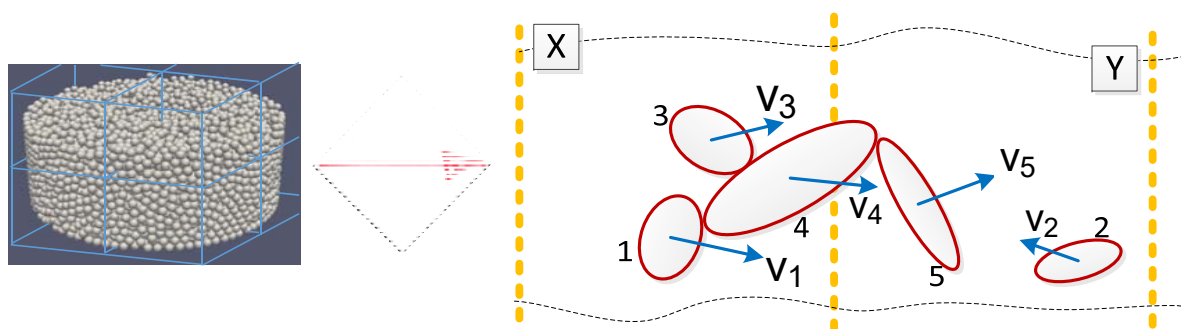


Computation Using Multiple CPUs & MPI



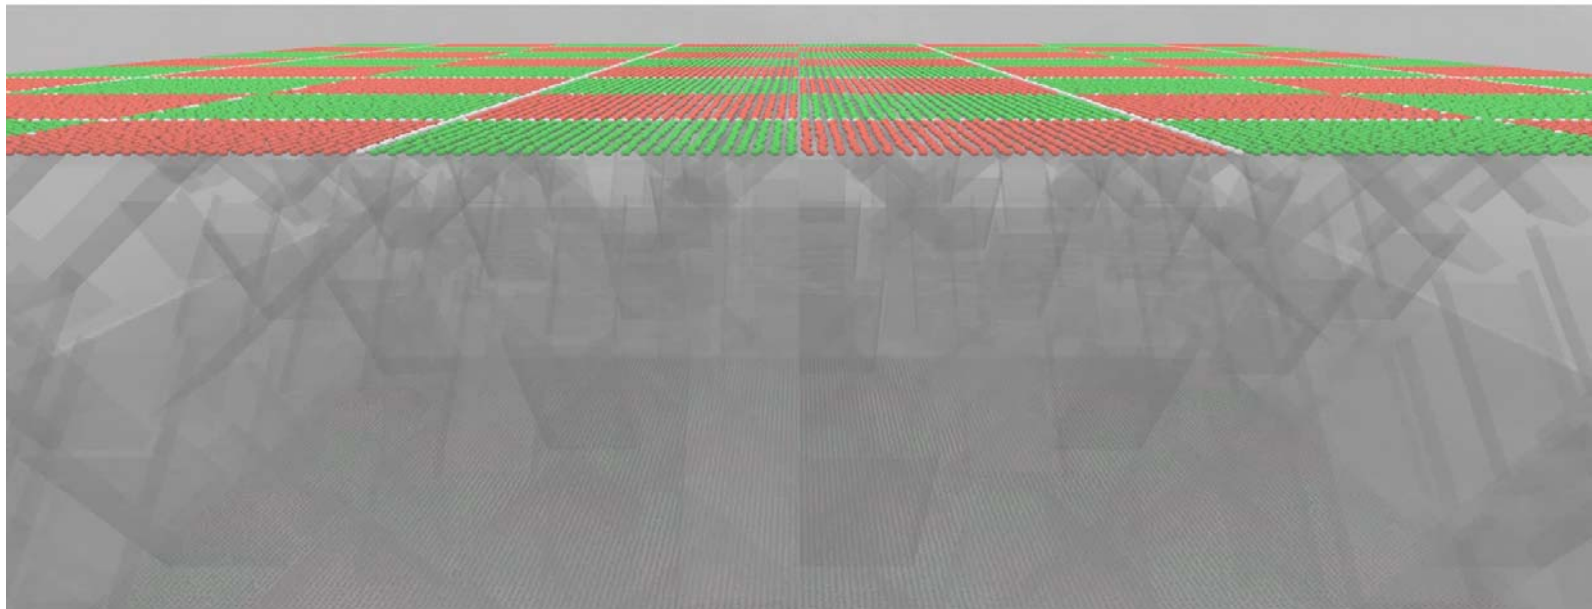
Breaking Up Dynamics for Distributed Computing

- Simulation divided into chunks executed on different cores
- Elements leave one chunk (subdomain) to move to a different one
- Key issues:
 - Dynamic load balancing
 - Establish a dynamic data exchange (DDE) protocol to implement migration at run time



0.5 Million Bodies on 64 Cores

[Penalty Approach, MPI-based]



Wrapping Things Up

- In the midst of a commoditization of parallel computing
 - Consequence of Moore's law
- Many alternatives, landscape very fluid
 - Lots of hardware options: GPU, CPU (Intel/AMD), Phi, clusters, FPGAs, etc.
 - Lots of software ecosystems: CUDA, OpenCL, OpenMP, OpenACC, MPI, Charm++, etc.
- Parallel computing here to stay
 - Can pursue new physics + higher fidelity + bridge space/time scales

Chrono Reliance on Parallel Computing

- GPU Computing
 - Granular material
 - Terrain modeling
 - Fluid-solid interaction (CFD)
- Multi-core parallel computing
 - Workhorse in Chrono::Parallel
 - Granular material
 - Terrain modeling
- Distributed memory parallel computing
 - Chrono::Vehicle
 - Pilot implementation, Chrono::FSI
 - Pilot implementation, terrain simulation
- GPU Computing + Multi-core
 - Pilot implementation, meshless methods for mud/silt

Looking Ahead, Chrono Reliance on Parallel Computing



- Distributed memory for terrain
- GPU computing for FSI (fording, sloshing, underwater robotics)